

# Lecture 5 — Interrupt Implementation

Jeff Zarnett & Mike Cooper-Stachowsky  
[jzarnett@uwaterloo.ca](mailto:jzarnett@uwaterloo.ca), [mstachowsky@uwaterloo.ca](mailto:mstachowsky@uwaterloo.ca)

Department of Electrical and Computer Engineering  
University of Waterloo

March 31, 2024

Credit: Based on material from Allyson Giannikouris

The previous topic discussed interrupts, both hardware and software.

We also learned about how to deal with them on the Cortex M4.

But to truly understand them, we need to look deeper into hardware.

Avoid polling as much as possible: it's a waste of time!

Instead: CPU can do whatever else is ready...

And get alerted when something happens.

**OH, YOU WANTED TO  
SLEEP IN TODAY?**



**LET ME PLAY YOU THE  
SONG OF MY PEOPLE.**

goldfunk.com

**Me: I need to go to sleep earlier**

**Also me:**



The CPU can be in a “sleeping” mode, but we’re ignoring that for this course.

What can cause an interrupt?

- Clock signals
- Timers
- Peripherals (UART, SPI, I2C, GPIO, ADC...)

We covered 3 of the 4 items here last time:

- 1** Interrupt the CPU
- 2** Save the state of the current task (PC, SP, R0-R3, LR)
- 3** Handle the interrupt
- 4** Resume execution (unless it's hard fault or similar)

But we didn't cover how the hardware interrupt works!

# How to get the CPU's Attention



It's a physical voltage signal (the power of electricity!).

# How to get the CPU's Attention

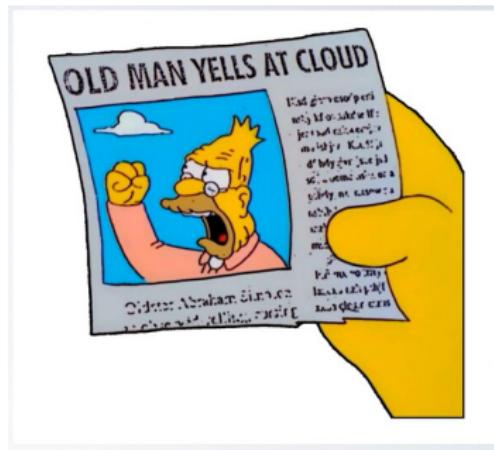
Each interrupt has a dedicated wire – we call them **lines**.

The line indicates the interrupt is ready.

The CPU has various interrupt registers to see if the interrupt is ready.

Signalling the processor is called an **interrupt request (IRQ)**.

Configuring the interrupts is a bit of a headache and at one time involved setting physical jumpers on cards to configure them.



But the idea is still a table of interrupts...

We use an array of function pointers!

Each interrupt references a particular address in the array.

Cortex M's vector table starts at 0x0.

# Interrupt Vector Table

Exception number	IRQ number	Offset	Vector
16+n	n	0x0040+4n	IRQn
.	.	.	.
.	.	.	.
18	2	0x004C	IRQ2
17	1	0x0048	IRQ1
16	0	0x0044	IRQ0
15	-1	0x0040	Systick
14	-2	0x003C	PendSV
13		0x0038	Reserved
12			Reserved for Debug
11	-5	0x002C	SVCall
10			
9			
8			Reserved
7			
6	-10	0x0018	Usage fault
5	-11	0x0014	Bus fault
4	-12	0x0010	Memory management fault
3	-13	0x0008	Hard fault
2	-14	0x000C	NMI
1		0x0008	Reset
		0x0004	Initial SP value
		0x0000	

Vector tables are configured in the setup code.

Good news: we almost never have to write it ourselves.

Pre-made code names the interrupt handler functions...  
And the names are used to write your IRQ functions.

## Setting up the vector table

```
.section .isr_vector,"a",%progbits
.type g_pfnVectors, %object
.size g_pfnVectors, .-g_pfnVectors

g_pfnVectors:
.word _estack
.word Reset_Handler
.word NMI_Handler
.word HardFault_Handler
.word MemManage_Handler
.word BusFault_Handler
.word UsageFault_Handler
.word 0
.word 0
.word 0
.word 0
.word SVC_Handler
.word DebugMon_Handler
.word 0
.word PendSV_Handler
.word SysTick_Handler
```

## A sample, pre-made Interrupt Service Routine (ISR)

```
.weak NMI_Handler
.thumb_set NMI_Handler,Default_Handler

.weak HardFault_Handler
.thumb_set HardFault_Handler,Default_Handler

.weak MemManage_Handler
.thumb_set MemManage_Handler,Default_Handler

.weak BusFault_Handler
.thumb_set BusFault_Handler,Default_Handler

.weak UsageFault_Handler
.thumb_set UsageFault_Handler,Default_Handler

.weak SVC_Handler
.thumb_set SVC_Handler,Default_Handler
```



This doesn't apply to SVC, though.

The code below runs when the SysTick interrupt happens, only if enabled.

---

```
void SysTick_Handler() {
    // Do the thing
}
```

---

Is it magical auto-generated design? Sure is.

Four main reasons to keep **Interrupt Subroutines (ISRs)** short:

- 1 Want to return to normal operation quickly
- 2 Minimize the chance of other interrupts
- 3 Shorter functions = less code = less risk of bug
- 4 C is hard enough; Assembly is much harder

It's better to avoid calling functions in ISRs.

We'll learn more about why this is complex in a future topic.

Functions called must be **reentrant**.

Any function that accesses a global or static variable is non-reentrant.

---

```
int tmp;
void swap( int *x, int *y ) {
    tmp = *x;
    *x = *y;
    *y = tmp;
}
```

---

How would you fix this to make it reentrant?

To return from an interrupt there's a special instruction `rti`.  
You cannot just do a regular `return` statement.

Let's think for a minute - why not?

Usually the IDE helps you out here.

# Stop Talking While I'm Interrupting

We know a few ways of handling interrupts during another interrupt:  
Ignore it, do it sequentially, switch to the new one?

But how to choose... by priority perhaps?

Interrupts are assigned a priority number; lower number is higher priority.  
At least in the Cortex M4; other systems are different.

# Stop Talking While I'm Interrupting

Some interrupts are always at a high priority and we can't change that.  
Others are under our control!

We aren't going to handle nested interrupts in our OS (too hard).

Some interrupts are **maskable**: can be ignored.



IRQs fall in this category. Usually indicate non-critical things.

Other interrupts are non-maskable (NMI).



Cannot be turned off or ignored; for critical things.

These force the CPU to halt execution!

In the Cortex M4 there's only one NMI.

It's the one that gets called when other interrupt handlers fault.

It's higher priority than anything other than RESET.

# Fault Priority Levels

Type	Exception Number	IRQ Number	Exception Type	Priority	Handled Using
processor core exceptions or internal interrupts	1	n/a	Reset	-3	
	2	-14	NMI	-2	
	3	-13	HardFault	-1	
	4	-12	MemManage	Configurable	Fault handlers
	5	-11	BusFault		
	6	-10	UsageFault		
	11	-5	SVCall		
	14	-2	PendSV		System handlers
	15	-1	SYSTic		
device-specific exceptions or external interrupts	16	0	IRQ		
	17	1			
	...	...			ISRs

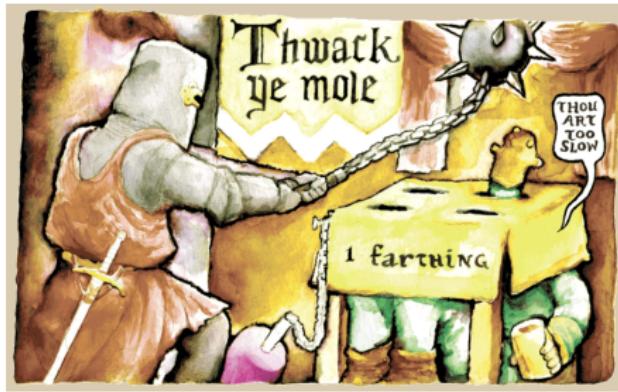
**Pending:** The conditions for the interrupt are met, but ISR has not run.

**Active:** The CPU is running the ISR.

**Pending and Active:** The next one is ready but we aren't done with the last!

**Inactive:** Neither pending nor active.

Pending and active can happen if we're not keeping up with the speed of the device somehow, or the device isn't keeping up with the CPU.



Sometimes it's also about luck and timing; events can happen very close together.

There is a single setting to turn on/off all interrupts.

When “off”, if triggered, they go to the pending state.

When interrupts re-enabled, all pending ones evaluated by priority.

This does not turn off NMI.

When a hardware interrupt occurs, though, what does the ISR do?

Simple ISRs are easier to write but they can only do so much.

What if the interrupt handler just needs to tell a user program it can proceed?

Don't forget to Like and Subscribe

This takes us down an interesting path...



Tasks (user processes) have state, more than just the current registers.

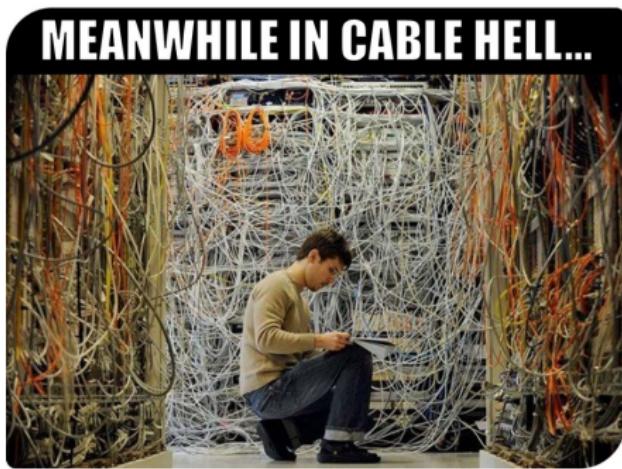
We'll have to return to the subject of how tasks (processes and threads) can be made to wait when they're not able to continue.

Just to be clear, what we're talking about is context-switching and multi-threading, but we're not ready for that yet.

It will come up soon and will be important in the labs.

# On to Hardware Devices Then

Before we're ready to go on, we need to spend some more time on hardware.



Our next topic will focus on communicating with hardware devices!